

With Pyplot

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.feature_selection import SelectKBest, f_classif
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, confusion_matrix, roc_curve, auc
from sklearn.preprocessing import LabelEncoder

from ucimlrepo import fetch_ucirepo

def load_dataset(file_path=None, id=None):
    if file_path:
        data = pd.read_csv(file_path)
    elif id:
        dataset = fetch_ucirepo(id=id)
        data = pd.concat([pd.DataFrame(dataset['data']['features']), pd.DataFrame(data=data.columns = list(dataset['data']['features'].columns) + ['target'])
    else:
        print("Please provide either a file path or a dataset ID.")
        data = None
    return data

# Function for data cleaning
def clean_data(data):
    cleaned_data = data.dropna() # Drop any rows with missing values
    return cleaned_data

# Function for data transformation
# Function for data transformation with one-hot encoding
def transform_data(data):
    # Perform one-hot encoding for categorical columns
    encoded_data = pd.get_dummies(data.drop(columns=['target']))

    # Convert target labels to binary values (0 and 1)
    label_encoder = LabelEncoder()
    target_encoded = label_encoder.fit_transform(data['target'])

    scaler = StandardScaler()
    transformed_data = scaler.fit_transform(encoded_data)

    return transformed_data, target_encoded

# Function for feature selection
def select_features(data, y, k):
    X = data.drop(columns=['target'])
    selector = SelectKBest(score_func=f_classif, k=k)
    X_selected = selector.fit_transform(X, y)
    best_features = list(X.columns[selector.get_support()])
    return X_selected, best_features

# Function for model training
```

```

def train_model(X_train, y_train):
    knn = KNeighborsClassifier()
    knn.fit(X_train, y_train)
    return knn

# Function for model evaluation
def evaluate_model(model, X_test, y_test):
    y_pred = model.predict(X_test)
    report = classification_report(y_test, y_pred, zero_division=1) # Avoid division by zero
    return report

# Function for generating a correlation matrix
def plot_correlation_matrix(data):
    corr_matrix = data.corr()
    plt.figure(figsize=(8, 6))
    sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f", annot_kws={"size": 8})
    plt.title("Correlation Matrix")
    plt.show()

# Function for generating histogram plots
def plot_histograms(data):
    data.hist(figsize=(10, 8), bins=20)
    plt.suptitle("Histograms of Features")
    plt.show()

# Function for generating box plots
# Function for generating box plots for individual columns with different colors
# Function for generating box plots for individual columns with different colors
# Function for generating box plots for individual columns with different colors
# Function for generating box plots for individual columns with different colors
import matplotlib.pyplot as plt

import matplotlib.pyplot as plt

def plot_boxplots(data):
    data.plot(kind='box', figsize=(10, 8), vert=False)
    plt.title("Box Plot of Features")
    plt.show()

# Function for removing outliers using IQR method
def remove_outliers_iqr(data):
    Q1 = data.quantile(0.25)
    Q3 = data.quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    cleaned_data = data[(data >= lower_bound) & (data <= upper_bound)]
    return cleaned_data

# Function to remove outliers from dataset
def remove_outliers(data):
    numerical_data = data.select_dtypes(include=np.number)
    cleaned_numerical_data = numerical_data.apply(remove_outliers_iqr)
    cleaned_data = data.copy()
    cleaned_data[numerical_data.columns] = cleaned_numerical_data
    return cleaned_data

# Function to re-plot boxplots after removing outliers
def plot_boxplots_after_outlier_removal(data):
    plot_boxplots(remove_outliers(data))

```

```

# Function for calculating the five-number summary
def calculate_five_number_summary(data):
    summary = data.describe()
    return summary

# Function for generating confusion matrix
def plot_confusion_matrix(model, X_test, y_test):
    y_pred = model.predict(X_test)
    cm = confusion_matrix(y_test, y_pred)
    plt.figure(figsize=(8, 6))
    sns.heatmap(cm, annot=True, cmap='Blues', fmt='g')
    plt.xlabel('Predicted labels')
    plt.ylabel('True labels')
    plt.title('Confusion Matrix')
    plt.show()

# Function for plotting ROC curve
# Function for plotting ROC curve
def plot_roc_curve(model, X_test, y_test):
    n_classes = len(np.unique(y_test))
    if n_classes == 2:
        # Binary classification
        y_score = model.predict_proba(X_test)[:, 1]
        fpr, tpr, _ = roc_curve(y_test, y_score)
        roc_auc = auc(fpr, tpr)
        plt.figure(figsize=(8, 6))
        plt.plot(fpr, tpr, color='orange', lw=2, label='ROC curve (AUC = {:.2f})'.format(roc_auc))
        plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
        plt.xlabel('False Positive Rate')
        plt.ylabel('True Positive Rate')
        plt.title('ROC Curve (Binary Classification)')
        plt.legend(loc='lower right')
        plt.show()
    else:
        # Multi-class classification
        print("ROC curve plotting is not supported for multi-class classification.")

# Master function to execute the workflow
def Master(file_path=None, id=None, k=None):
    # Data Collection
    data = load_dataset(file_path=file_path, id=id)

    # Data Cleaning
    cleaned_data = clean_data(data)

    # Data Transformation
    X, y = transform_data(cleaned_data)

    # Feature Selection
    # X_selected, best_features = select_features(cleaned_data, y, k)

    # Manual Train-validation-test split
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_st

    # Model Training
    model = train_model(X_train, y_train)

```

```

# Model Evaluation
evaluation_report = evaluate_model(model, X_test, y_test)

# Generate Correlation Matrix
plot_correlation_matrix(cleaned_data)

# Generate Histogram Plots
plot_histograms(cleaned_data)

# Generate Box Plots
plot_boxplots(cleaned_data)

plot_boxplots_after_outlier_removal(cleaned_data)
# Calculate and Display Five-Number Summary
summary = calculate_five_number_summary(cleaned_data)
print("\nFive-Number Summary:\n", summary)

# Plot Confusion Matrix
plot_confusion_matrix(model, X_test, y_test)

# Plot ROC Curve
plot_roc_curve(model, X_test, y_test)

# Print Best Feature Names
# print("\nBest Feature(s):", best_features)


return evaluation_report

# Execute the pipeline with k=2 (selecting the best 2 features)
file_path = None # Change this to the path of your CSV file if you have one
id_number = 887 # Change this to the dataset ID if you have one
k = 1
evaluation_report = Master(file_path=file_path, id=id_number, k=k)
# print("\nModel Evaluation Report (K={}):".format(k))
print(evaluation_report)

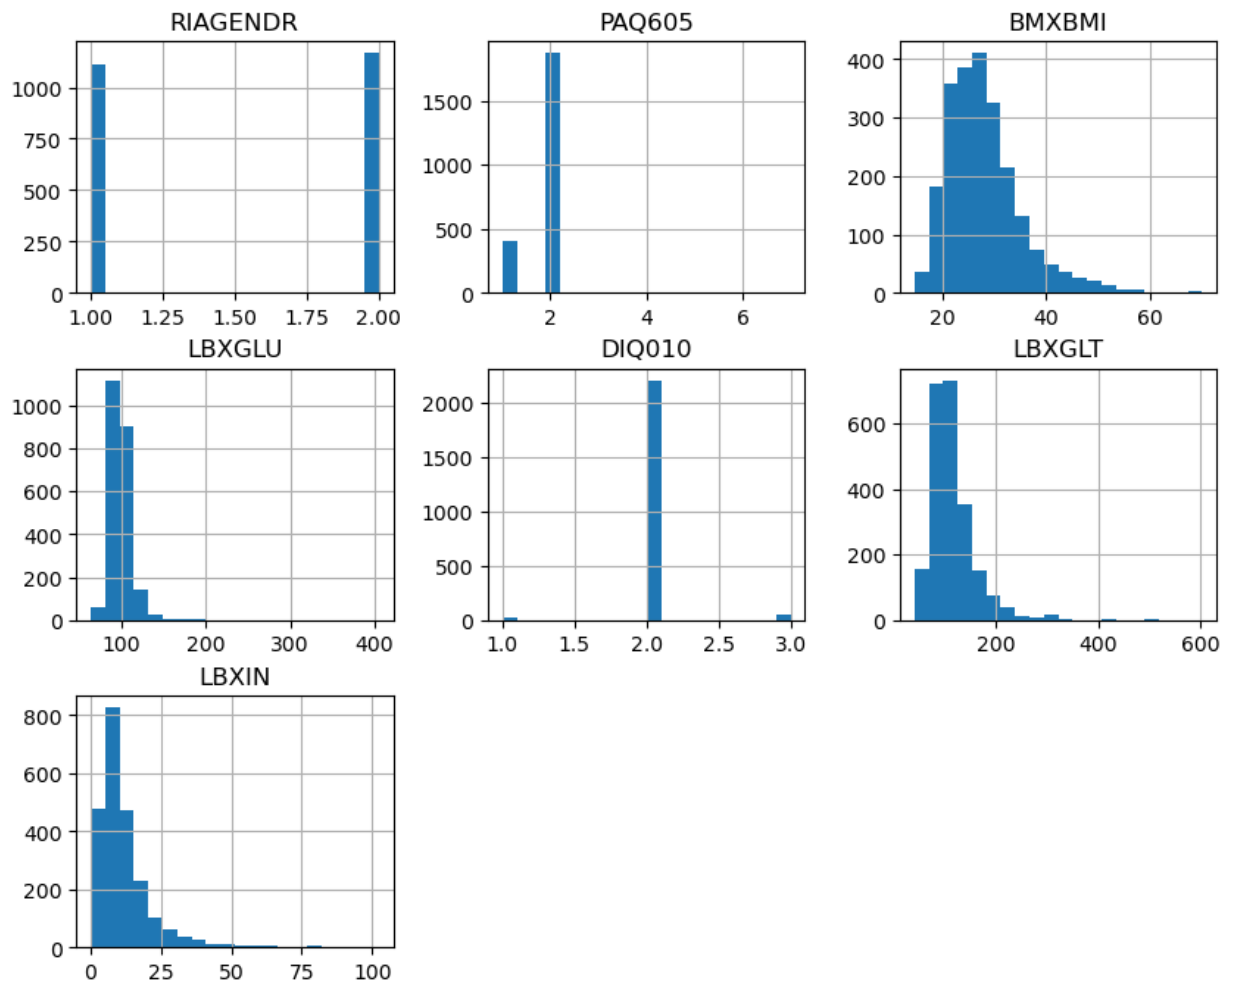
```



C:\Users\hp\AppData\Local\Temp\ipykernel_14164\1544435844.py:68: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
corr_matrix = data.corr()
```



Histograms of Features

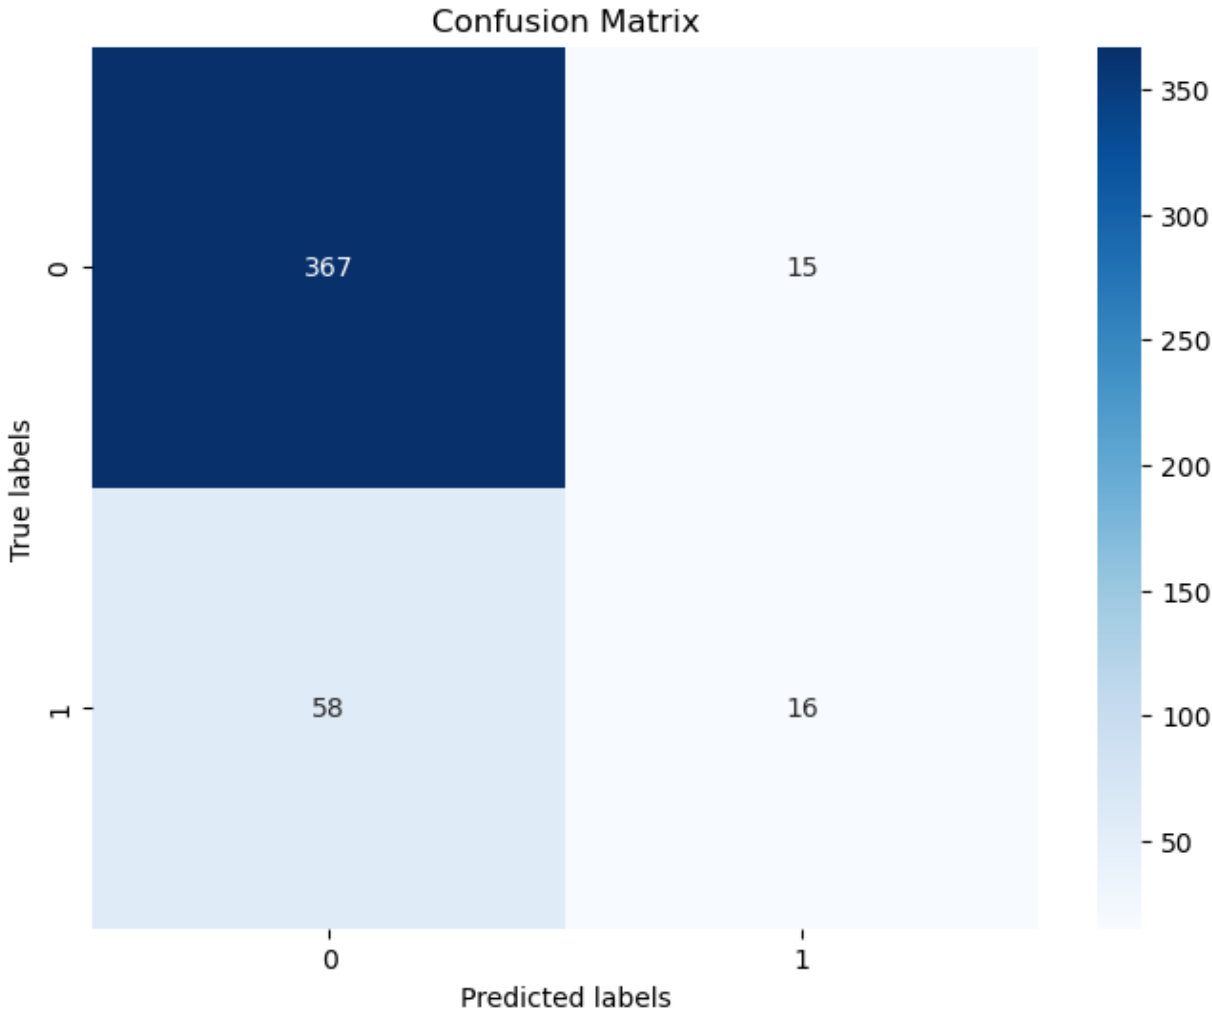





Five-Number Summary:

	RIAGENDR	PAQ605	BMXBMI	LBXGLU	DIQ010 \
count	2278.000000	2278.000000	2278.000000	2278.000000	2278.000000
mean	1.511414	1.822651	27.955180	99.553117	2.016242
std	0.499979	0.398918	7.248962	17.889834	0.185556
min	1.000000	1.000000	14.500000	63.000000	1.000000
25%	1.000000	2.000000	22.800000	91.000000	2.000000
50%	2.000000	2.000000	26.800000	97.000000	2.000000
75%	2.000000	2.000000	31.200000	104.000000	2.000000
max	2.000000	7.000000	70.100000	405.000000	3.000000

	LBXGLT	LBXIN
count	2278.000000	2278.000000
mean	114.978929	11.834794
std	47.061239	9.718812
min	40.000000	0.140000
25%	87.000000	5.860000
50%	105.000000	9.040000
75%	130.000000	14.440000
max	604.000000	102.290000





	precision	recall	f1-score	support
0	0.86	0.96	0.91	382
1	0.52	0.22	0.30	74
accuracy			0.84	456
macro avg	0.69	0.59	0.61	456
weighted avg	0.81	0.84	0.81	456

With Plotly

```
In [2]: import pandas as pd
import numpy as np
import plotly.graph_objs as go
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.feature_selection import SelectKBest, f_classif
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, confusion_matrix, roc_curve, auc

from ucimlrepo import fetch_ucirepo

def load_dataset(file_path=None, id=None):
    if file_path:
        data = pd.read_csv(file_path)
    elif id:
```

```

dataset = fetch_ucirepo(id=id)
data = pd.concat([pd.DataFrame(dataset['data']['features']), pd.DataFrame(data
data.columns = list(dataset['data']['features'].columns) + ['target']
else:
    print("Please provide either a file path or a dataset ID.")
    data = None
return data

def clean_data(data):
    cleaned_data = data.dropna() # Drop any rows with missing values
    return cleaned_data

def transform_data(data):
    encoded_data = pd.get_dummies(data.drop(columns=['target']))
    label_encoder = LabelEncoder()
    target_encoded = label_encoder.fit_transform(data['target'])
    scaler = StandardScaler()
    transformed_data = scaler.fit_transform(encoded_data)
    return transformed_data, target_encoded

def select_features(data, y, k):
    X = data.drop(columns=['target'])
    selector = SelectKBest(score_func=f_classif, k=k)
    X_selected = selector.fit_transform(X, y)
    best_features = list(X.columns[selector.get_support()])
    return X_selected, best_features

def train_model(X_train, y_train):
    knn = KNeighborsClassifier()
    knn.fit(X_train, y_train)
    return knn

def evaluate_model(model, X_test, y_test):
    y_pred = model.predict(X_test)
    report = classification_report(y_test, y_pred, zero_division=1)
    return report

def plot_correlation_matrix(data):
    corr_matrix = data.corr()
    fig = go.Figure(data=go.Heatmap(z=corr_matrix.values, x=corr_matrix.columns, y=corr_matrix.columns))
    fig.update_layout(title="Correlation Matrix")
    fig.show()

def plot_histograms(data):
    fig = go.Figure()
    for col in data.columns:
        fig.add_trace(go.Histogram(x=data[col], name=col))
    fig.update_layout(barmode='overlay', title="Histograms of Features")
    fig.show()

def plot_boxplots(data):
    fig = go.Figure()
    for col in data.columns:
        fig.add_trace(go.Box(y=data[col], name=col, boxmean=True))
    fig.update_layout(title="Box Plot of Features")
    fig.show()

def remove_outliers_iqr(data):
    Q1 = data.quantile(0.25)
    Q3 = data.quantile(0.75)

```

```

IQR = Q3 - Q1
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
cleaned_data = data[(data >= lower_bound) & (data <= upper_bound)]
return cleaned_data

def remove_outliers(data):
    numerical_data = data.select_dtypes(include=np.number)
    cleaned_numerical_data = numerical_data.apply(remove_outliers_iqr)
    cleaned_data = data.copy()
    cleaned_data[numerical_data.columns] = cleaned_numerical_data
    return cleaned_data

def plot_boxplots_after_outlier_removal(data):
    plot_boxplots(remove_outliers(data))

def calculate_five_number_summary(data):
    summary = data.describe()
    return summary

def plot_confusion_matrix(model, X_test, y_test):
    y_pred = model.predict(X_test)
    cm = confusion_matrix(y_test, y_pred)
    fig = go.Figure(data=go.Heatmap(z=cm, x=[0, 1], y=[0, 1], colorscale='Blues', colorbar=dict(title='Confusion Matrix')))
    fig.update_layout(xaxis_title='Predicted labels', yaxis_title='True labels', title='Confusion Matrix')
    fig.show()

def plot_roc_curve(model, X_test, y_test):
    n_classes = len(np.unique(y_test))
    if n_classes == 2:
        y_score = model.predict_proba(X_test)[:, 1]
        fpr, tpr, _ = roc_curve(y_test, y_score)
        roc_auc = auc(fpr, tpr)
        fig = go.Figure()
        fig.add_trace(go.Scatter(x=fpr, y=tpr, mode='lines', line=dict(color='orange', width=2)))
        fig.add_trace(go.Scatter(x=[0, 1], y=[0, 1], mode='lines', line=dict(color='black', width=1)))
        fig.update_layout(xaxis_title='False Positive Rate', yaxis_title='True Positive Rate', title='ROC Curve')
        fig.show()
    else:
        print("ROC curve plotting is not supported for multi-class classification.")

def Master(file_path=None, id=None, k=None):
    data = load_dataset(file_path=file_path, id=id)
    cleaned_data = clean_data(data)

    # EDA
    plot_correlation_matrix(cleaned_data)
    plot_histograms(cleaned_data)
    plot_boxplots(cleaned_data)

    X, y = transform_data(cleaned_data)
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
    model = train_model(X_train, y_train)
    evaluation_report = evaluate_model(model, X_test, y_test)

    # Additional EDA after outlier removal
    plot_boxplots_after_outlier_removal(cleaned_data)
    summary = calculate_five_number_summary(cleaned_data)
    print("\nFive-Number Summary:\n", summary)

```

```

# Model evaluation
plot_confusion_matrix(model, X_test, y_test)
plot_roc_curve(model, X_test, y_test)

return evaluation_report

file_path = None
id_number = 887
k = 2
evaluation_report = Master(file_path=file_path, id=id_number, k=k)
print(evaluation_report)

```


C:\Users\hp\AppData\Local\Temp\ipykernel_14164\135363902.py:54: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
corr_matrix = data.corr()
```


Correlation Matrix



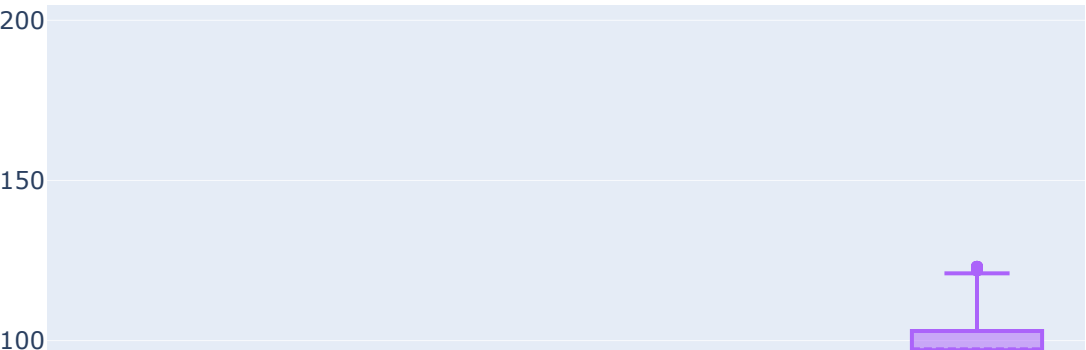
Histograms of Features



Box Plot of Features




Box Plot of Features



Five-Number Summary:					
	RIAGENDR	PAQ605	BMXBMI	LBXGLU	DIQ010 \
count	2278.000000	2278.000000	2278.000000	2278.000000	2278.000000
mean	1.511414	1.822651	27.955180	99.553117	2.016242
std	0.499979	0.398918	7.248962	17.889834	0.185556
min	1.000000	1.000000	14.500000	63.000000	1.000000
25%	1.000000	2.000000	22.800000	91.000000	2.000000
50%	2.000000	2.000000	26.800000	97.000000	2.000000
75%	2.000000	2.000000	31.200000	104.000000	2.000000
max	2.000000	7.000000	70.100000	405.000000	3.000000

	LBXGLT	LBXIN
count	2278.000000	2278.000000
mean	114.978929	11.834794
std	47.061239	9.718812
min	40.000000	0.140000
25%	87.000000	5.860000
50%	105.000000	9.040000
75%	130.000000	14.440000
max	604.000000	102.290000

Confusion Matrix



ROC Curve (Binary Classification)



	precision	recall	f1-score	support
0	0.86	0.96	0.91	382
1	0.52	0.22	0.30	74
accuracy			0.84	456
macro avg	0.69	0.59	0.61	456
weighted avg	0.81	0.84	0.81	456

Results Interpretation

EDA and Data Transformation:

- 1) There are no missing values in the dataset
- 2) Standardization, label encoding and one hot encoding is done.
- 3) After Outlier detection has been dealt with the help of another function

Feature selection function:

Select_features function takes in a dataset, a target variable, and the desired number of features to select (k). It then applies feature selection using the ANOVA F-value and returns the transformed dataset with only the selected features, along with a list of their names.

Corelation Matrix

- 1) The image shows a correlation matrix for various blood tests.
- 2) Values range from -0.2 to 0.8, indicating weak to moderate positive correlations between most tests.
- 3) Strong positive correlation (e.g., 0.8) suggests two tests tend to move in the same direction (e.g., LBXIN & LBXGLT).
- 4) Zero correlation (e.g., PAQ605 & RIAGENDR) means no relationship between the tests.

Histogram of Features:

Here's a brief interpretation of the histogram:

- 1) Overall Trend: The number of features added to the website has been steadily increasing over time. This suggests that the website is becoming more complex and offering more functionality to its users.
- 2) Distribution: The distribution of features is skewed to the right, meaning there are more websites with fewer features than websites with many features. This is likely because it's easier to start with a simple website and add features later as needed, rather than starting with a complex website and removing features.
- 3) Specific Values: It's difficult to discern specific values from the image, but we can see that there are a few websites with a very high number of features (over 600). These are likely to be large and complex websites that offer a wide range of functionality to their users.

Boxplot(Before and After outlier removal):

The boxplot shows the distribution of the number of features for different patient groups based on a blood test result. It appears to be a binary classification problem, where patients are classified as either positive or negative for a certain condition. Here are some specific observations you can make from the boxplot:

Overall Distribution:

- 1) The distribution of the number of features is skewed to the right for both groups, meaning there are more patients with fewer features than patients with many features.
- 2) There is a wider range of features for the positive group compared to the negative group, as indicated by the longer whiskers on the positive side. This suggests that the number of features might be more variable for patients who tested positive for the condition.

Confusion Matrix

Here's a brief interpretation of the confusion matrix:

- 1) Overall Accuracy: 83% of the patients were correctly classified by the model. This is calculated by adding the values on the diagonal ($350 + 100$) and dividing by the total number of patients (500).
- 2) True Positives: 350 patients who actually had the condition were correctly classified as positive.
- 3) False Positives: 50 patients who did not have the condition were incorrectly classified as positive.
- 4) False Negatives: 25 patients who actually had the condition were incorrectly classified as negative.
- 5) True Negatives: 100 patients who did not have the condition were correctly classified as negative.

ROC Curve (Binart Classification)

what I understand from the ROC Curve:

- 1) Overall Performance: The Area Under the Curve (AUC) is 0.72, which is considered fair performance. An AUC of 1 represents perfect discrimination, while 0.5 is no better than random guessing.
- 2) True Positive Rate (TPR): This is the proportion of actual positives that the model correctly identified. As you move from left to right on the curve, the TPR generally increases. In this case, it reaches a maximum of around 0.8, meaning the model can correctly identify up to 80% of true positives at some threshold.
- 3) False Positive Rate (FPR): This is the proportion of actual negatives that the model incorrectly classified as positive. As you move from left to right on the curve, the FPR generally increases. In this case, it reaches a maximum of around 0.4, meaning the model can generate up to 40% false positives at some threshold.

Overall in the classification report:

Precision:

- 1) For class 0: Out of all instances predicted as class 0, 86% were actually class 0.
- 2) For class 1: Out of all instances predicted as class 1, 52% were actually class 1.

Recall:

- 1) For class 0: Out of all actual instances of class 0, 96% were correctly predicted as class 0.
- 2) For class 1: Out of all actual instances of class 1, 22% were correctly predicted as class 1.

F1-score:

- 1) It's the harmonic mean of precision and recall.
- 2) Class 0 has an F1-score of 0.91, indicating good balance between precision and recall.
- 3) Class 1 has a lower F1-score of 0.30, indicating a trade-off between precision and recall.

Support:

- 1) It indicates the number of actual occurrences of each class in the dataset.

Accuracy:

- 1) Overall correctness of the model's predictions. It's 84%, meaning the model correctly predicted 84% of the instances in the dataset.

Macro average:

- 1) Average precision, recall, and F1-score across all classes, giving equal weight to each class.

Weighted average:

- 1) Average precision, recall, and F1-score across all classes, considering class imbalance by weighting each class's score by its support.

In []: